

kqkrpr8wf

February 15, 2025

### 0.0.1 Machine Learning Workflow for Auto Insurance Dataset

My dataset includes features related to **customer demographics, prior insurance history, claims data, premium adjustments, and conversion status**. Depending on my goal, different **ML techniques** will be applied.

---

## 0.1 1 Problem Definition

- I want to **predict whether a customer will convert** (buy insurance), it's a **classification problem**.
    - **Target Variable:** `Conversion_Status` (Binary: 0 = No, 1 = Yes)
    - **ML Techniques:** Logistic Regression, Random Forest, XGBoost, Neural Networks
  - I want to **predict premium adjustment** based on customer history, it's a **regression problem**.
    - **Target Variable:** `Premium_Adjustment_Credit` or `Premium_Adjustment_Region` (Continuous Values)
    - **ML Techniques:** Linear Regression, Random Forest Regressor, XGBoost Regressor
- 

## 0.2 2 Exploratory Data Analysis (EDA)

**Goal:** Understand dataset characteristics and relationships between variables.

**Techniques:**

- **Check for Missing Values:** `df.isnull().sum()`
  - **Feature Distributions:** Histograms, Boxplots
  - **Correlation Analysis:** `df.corr()` + Heatmap
  - **Categorical Data Visualization:** Countplots for `Marital_Status`, `Policy_Type`, etc.
  - **Outlier Detection:** IQR method
- 

## 0.3 3 Data Preprocessing & Feature Engineering

**Goal:** Prepare data for ML models by encoding categorical variables, handling missing values, and scaling numerical data.

#### Techniques:

##### - Encoding Categorical Variables:

- One-Hot Encoding (Marital\_Status, Region, Policy\_Type)
- Label Encoding (Claims\_Severity: Low → 0, Medium → 1, High → 2)

##### - Handling Missing Values:

- Impute NaN values with **mean/median** (numerical) or **mode** (categorical).

##### - Feature Scaling:

- **Standardization (StandardScaler)** for models like Logistic Regression, SVM
- **Normalization (MinMaxScaler)** for Neural Networks

##### - Feature Selection:

- Remove highly correlated features (e.g., Premium\_Adjustment\_Credit and Premium\_Adjustment\_Region may be related)
  - Use **Recursive Feature Elimination (RFE)** to find important features
- 

## 0.4 4 Model Selection & Training

Based on the problem type:

### 0.4.1 Classification (Predicting Conversion\_Status)

- **Logistic Regression** (Baseline Model)
- **Random Forest Classifier** (Handles categorical & numerical features well)
- **XGBoost/LightGBM** (Best for structured insurance data)
- **Artificial Neural Networks (ANN)** (For deep learning-based approaches)

#### Pipeline:

1. Train-Test Split: `train_test_split(df, test_size=0.2, stratify=df['Conversion_Status'])`
  2. Model Training: `RandomForestClassifier(n_estimators=100)`
  3. Evaluate with **Accuracy, Precision, Recall, F1-score, ROC-AUC**
- 

### 0.4.2 Regression (Predicting Premium\_Adjustment\_Credit)

- **Linear Regression** (Baseline)
- **Random Forest Regressor** (Handles nonlinear relationships)
- **XGBoost Regressor** (For boosting performance)

#### Pipeline:

1. Train-Test Split: `train_test_split(df, test_size=0.2, random_state=42)`
2. Model Training: `XGBRegressor(n_estimators=500, learning_rate=0.1)`
3. Evaluate using **RMSE, R<sup>2</sup> Score, MAE**

---

## 0.5 5 Model Evaluation & Optimization

### Classification Metrics:

- `classification_report(y_test, y_pred)` (Accuracy, Precision, Recall, F1-score)
- `roc_auc_score(y_test, y_pred_prob)` (For imbalanced data)

### Regression Metrics:

- `mean_squared_error(y_test, y_pred, squared=False)` (RMSE)
- `r2_score(y_test, y_pred)` (Explained Variance)

### Hyperparameter Tuning:

- `GridSearchCV` (For Logistic Regression, Random Forest)
  - `RandomizedSearchCV` (For XGBoost, LightGBM)
- 

## 0.6 6 Model Deployment (Optional)

If you want to **deploy the model**, I can:

- **Save the model:** `pickle.dump(model, open('model.pkl', 'wb'))`
  - **Deploy with Flask/FastAPI** for API-based predictions
  - **Build a Dashboard with Streamlit/Dash**
- 

### 0.6.1 Conclusion

My dataset allows for **multiple ML applications**, from customer conversion prediction to premium optimization. **Tree-based models (Random Forest, XGBoost)** will likely perform **best** due to structured data.

```
[16]: import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")
import pandas as pd
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import classification_report, accuracy_score, \
    confusion_matrix, roc_auc_score
from xgboost import XGBClassifier
from lightgbm import LGBMClassifier
```

```
from catboost import CatBoostClassifier
```

```
[ ]:
```

```
[ ]:
```

```
[3]: df = pd.read_csv("/kaggle/input/insurance-data-personal-auto-line-of-business/
↳synthetic_insurance_data.csv")
```

```
[4]: df.head()
```

```
[4]:   Age  Is_Senior  Marital_Status  Married_Premium_Discount  Prior_Insurance  \
0    47          0        Married                86        1-5 years
1    37          0        Married                86        1-5 years
2    49          0        Married                86        1-5 years
3    62          1        Married                86        >5 years
4    36          0         Single                0        >5 years
```

```
   Prior_Insurance_Premium_Adjustment  Claims_Frequency  Claims_Severity  \
0                                50                    0             Low
1                                50                    0             Low
2                                50                    1             Low
3                                0                     1             Low
4                                0                     2             Low
```

```
   Claims_Adjustment  Policy_Type  ...  Time_Since_First_Contact  \
0                   0  Full Coverage  ...                    10
1                   0  Full Coverage  ...                    22
2                   50  Full Coverage  ...                    28
3                   50  Full Coverage  ...                     4
4                  100  Full Coverage  ...                    14
```

```
   Conversion_Status  Website_Visits  Inquiries  Quotes_Requested  \
0                   0                5          1                 2
1                   0                5          1                 2
2                   0                4          4                 1
3                   1                6          2                 2
4                   1                8          4                 2
```


```
   Time_to_Conversion  Credit_Score  Premium_Adjustment_Credit  Region  \
0                   99           704                    -50  Suburban
1                   99           726                    -50    Urban
2                   99           772                    -50    Urban
3                    2           809                    -50    Urban
4                  10           662                     50  Suburban
```

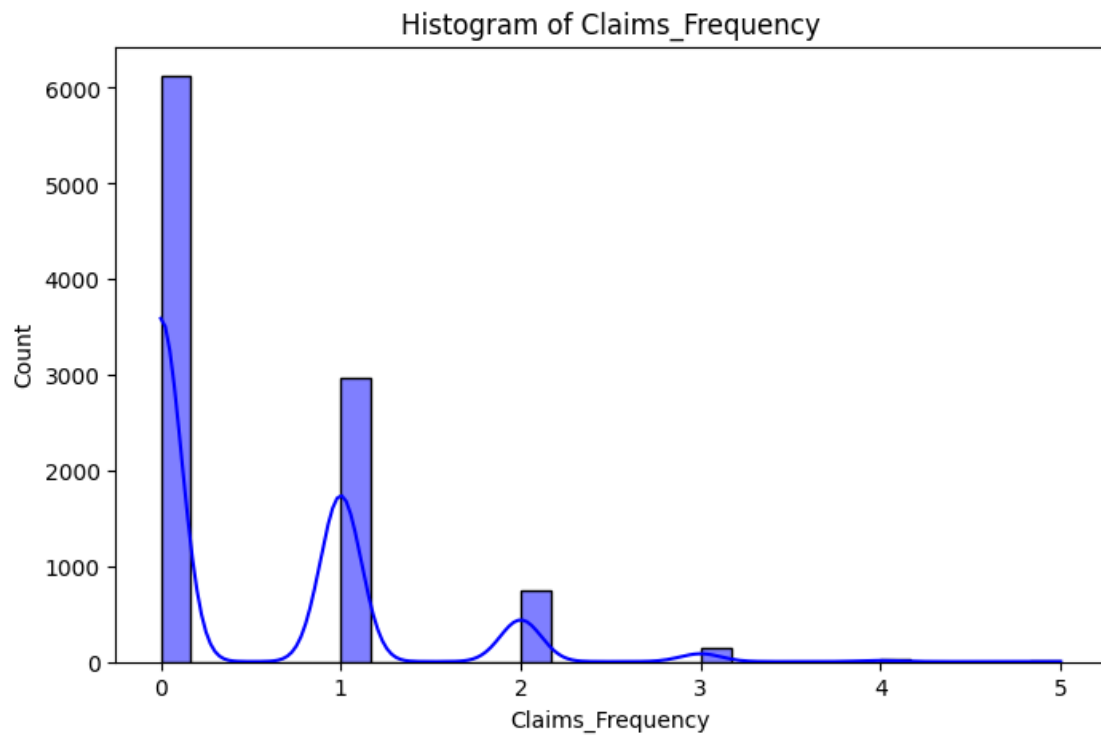
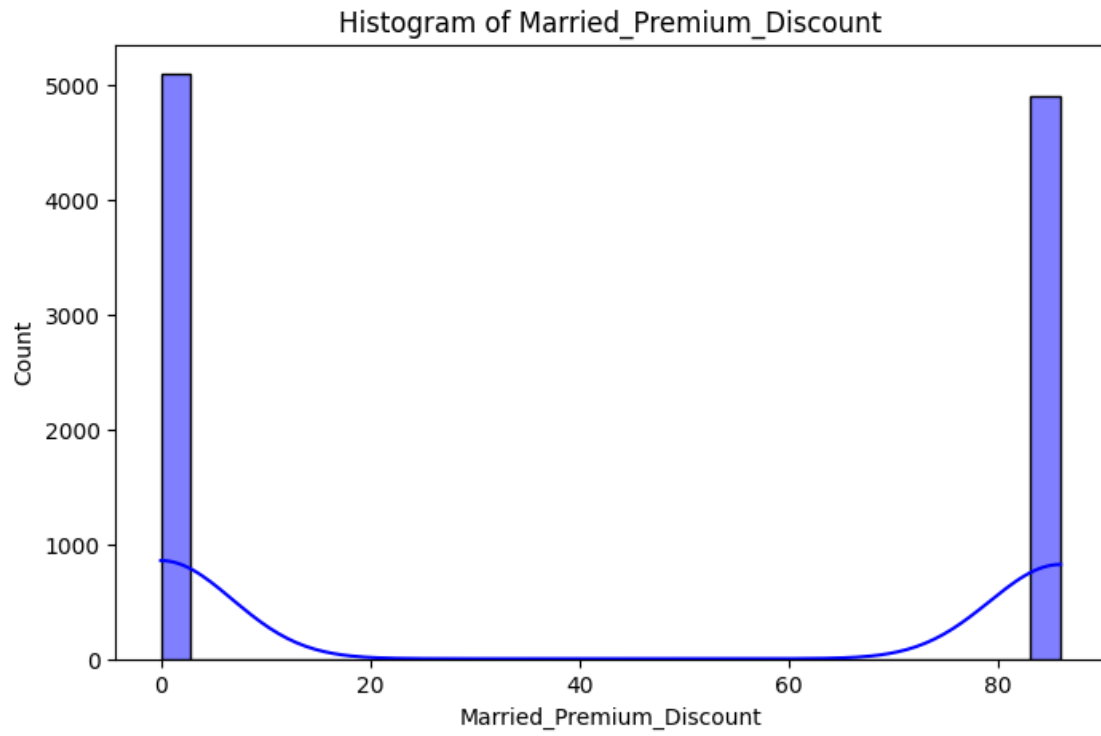
```
Premium_Adjustment_Region
```


0	50
1	100
2	100
3	100
4	50

[5 rows x 27 columns]

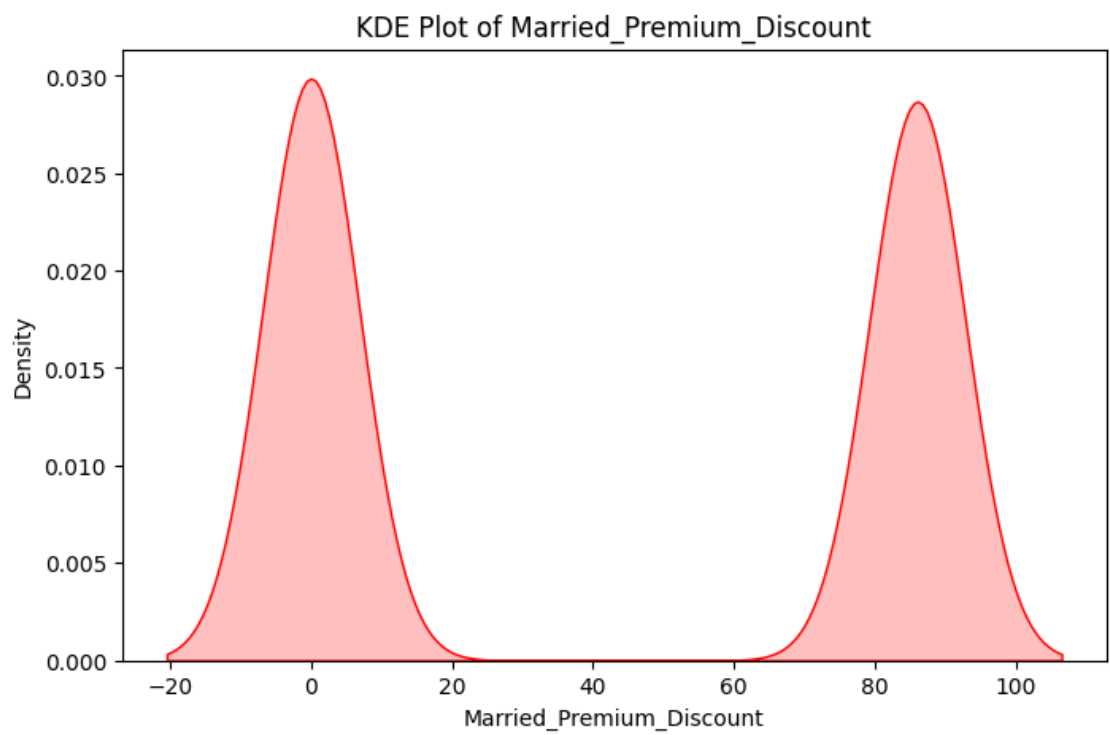

```
[5]: # Histogram
num_cols = ['Age', 'Married_Premium_Discount', 'Claims_Frequency', 'Credit_Score']
for col in num_cols:
    plt.figure(figsize=(8, 5))
    sns.histplot(df[col], bins=30, kde=True, color='blue')
    plt.title(f'Histogram of {col}')
    plt.show()
```




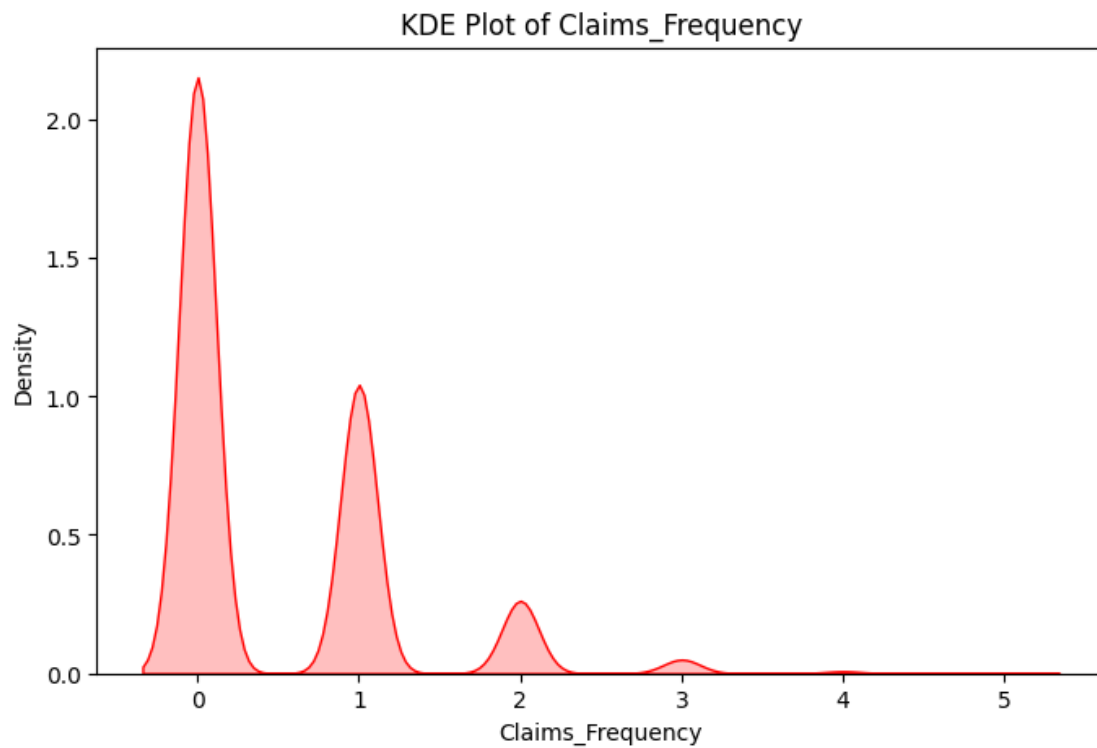





```
[6]: # KDE Plots
for col in num_cols:
    plt.figure(figsize=(8, 5))
    sns.kdeplot(df[col], fill=True, color='red')
    plt.title(f'KDE Plot of {col}')
    plt.show()
```

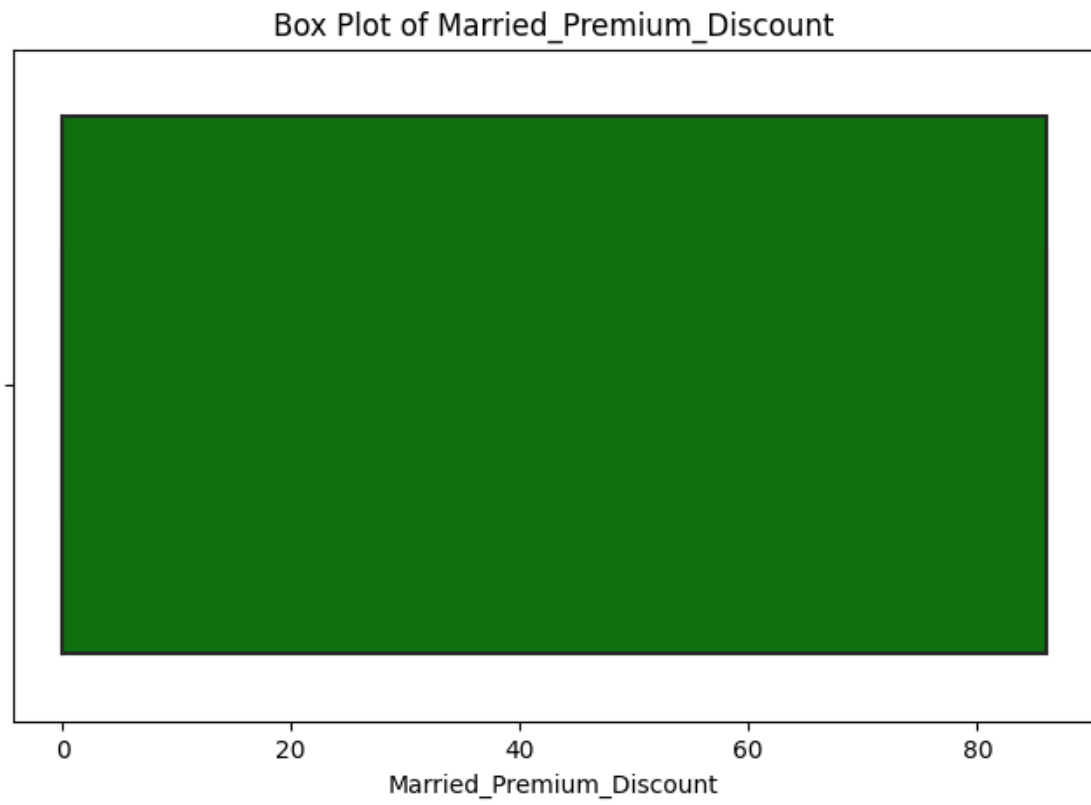


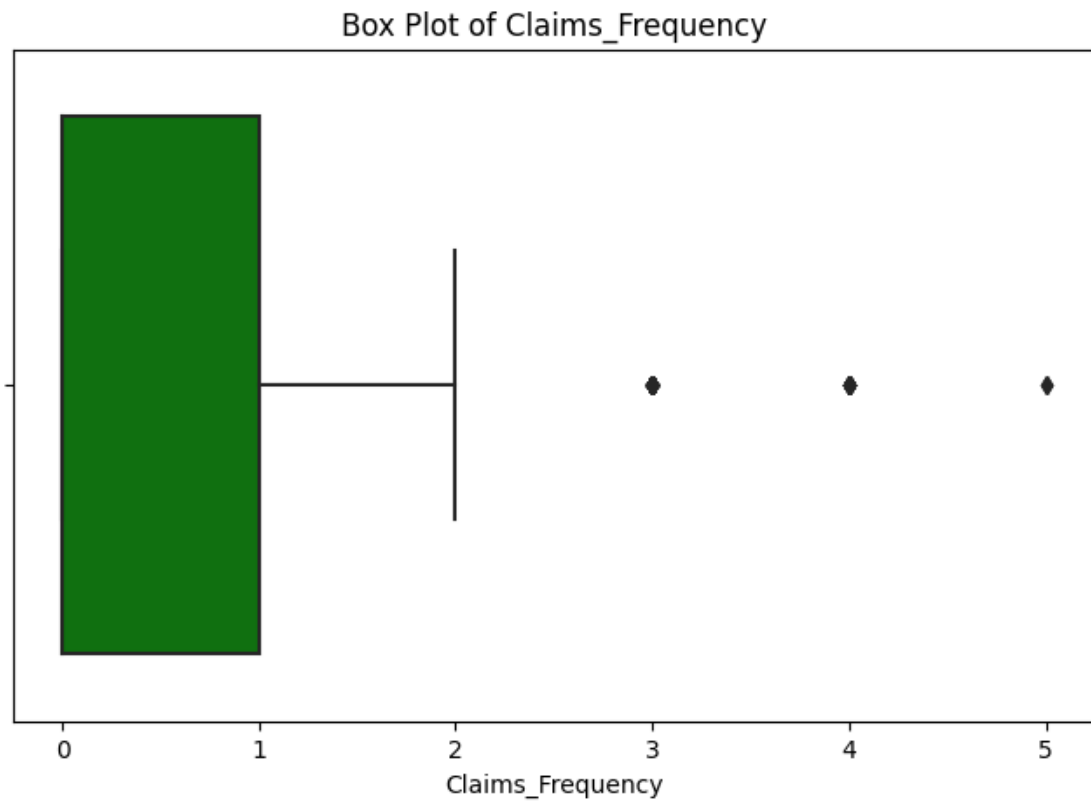





```
[7]: # Box Plots
for col in num_cols:
    plt.figure(figsize=(8, 5))
    sns.boxplot(x=df[col], color='green')
    plt.title(f'Box Plot of {col}')
    plt.show()
```




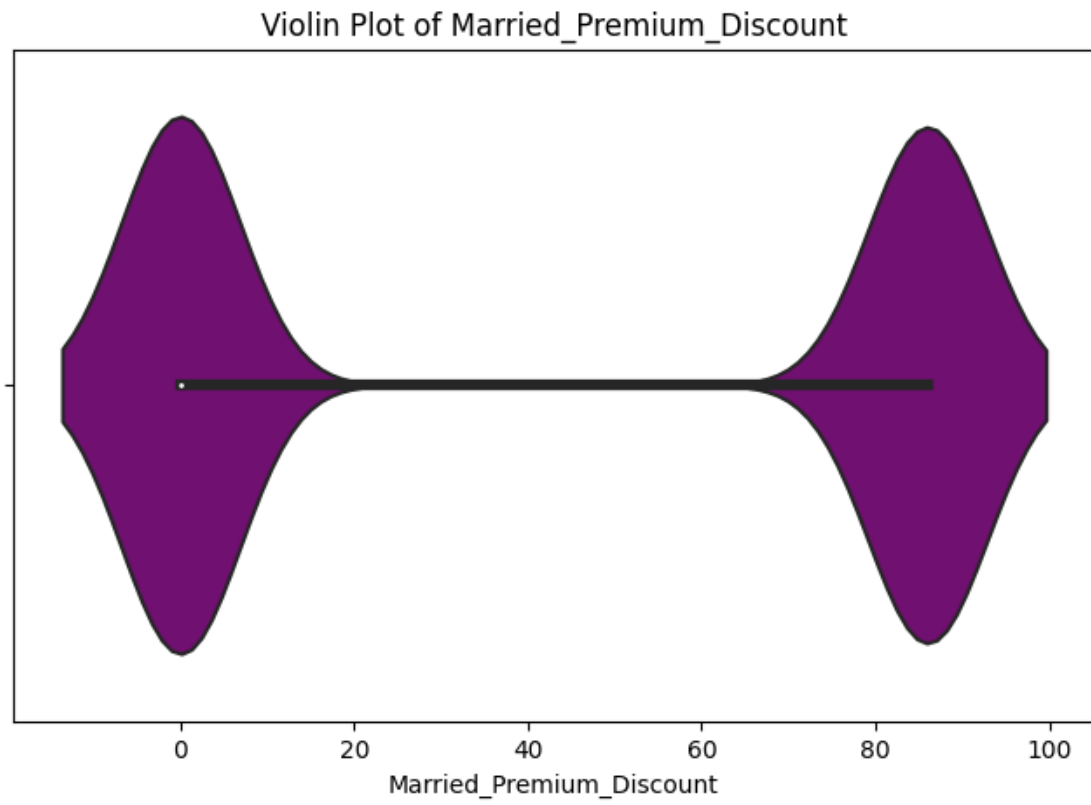


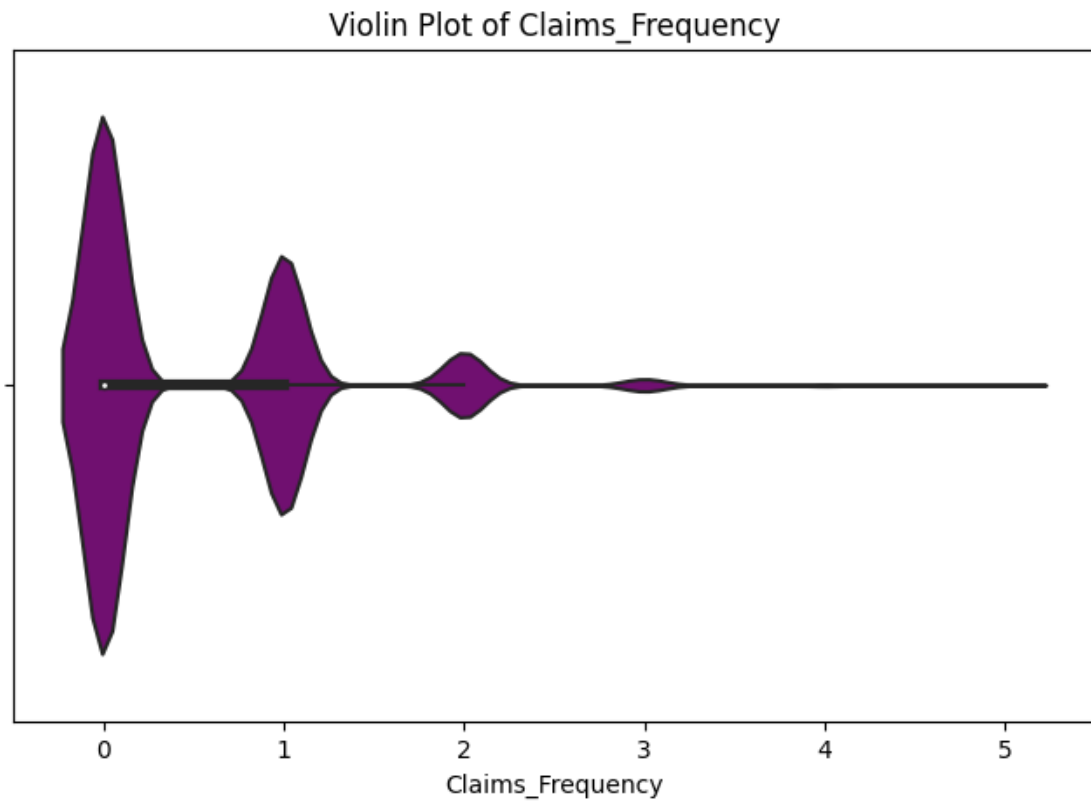




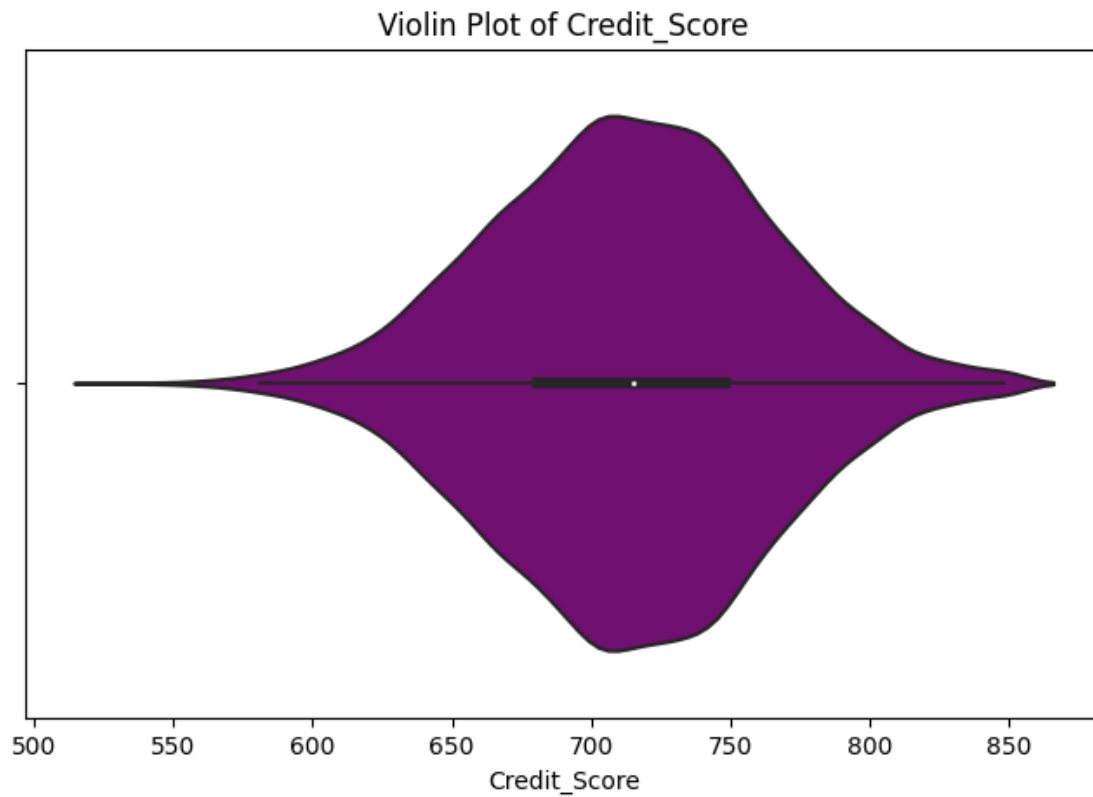
```
[8]: # Violin Plots
for col in num_cols:
    plt.figure(figsize=(8, 5))
    sns.violinplot(x=df[col], color='purple')
    plt.title(f'Violin Plot of {col}')
    plt.show()
```



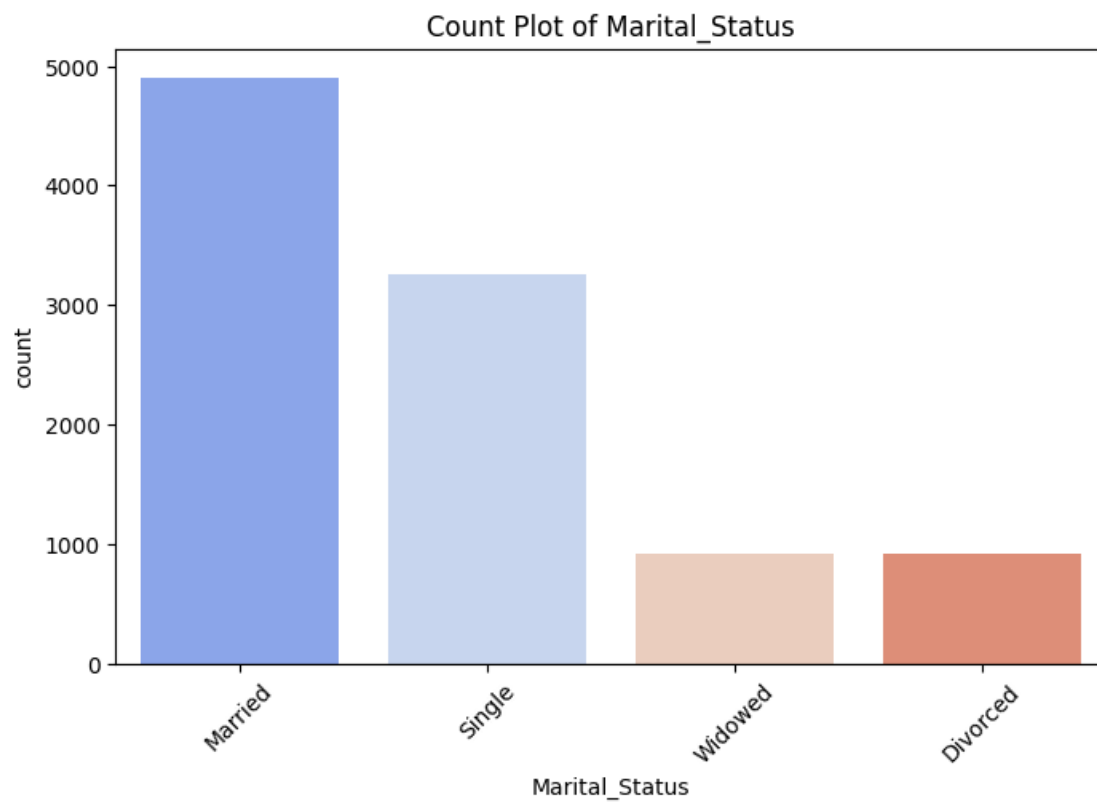


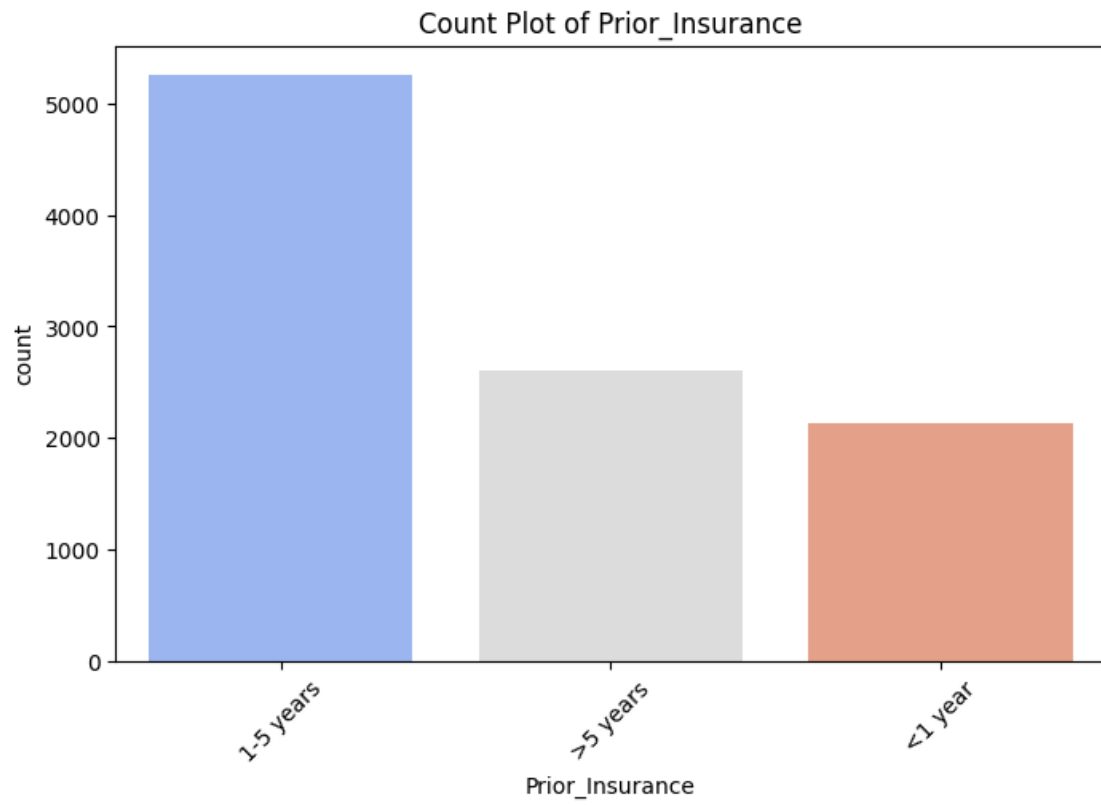


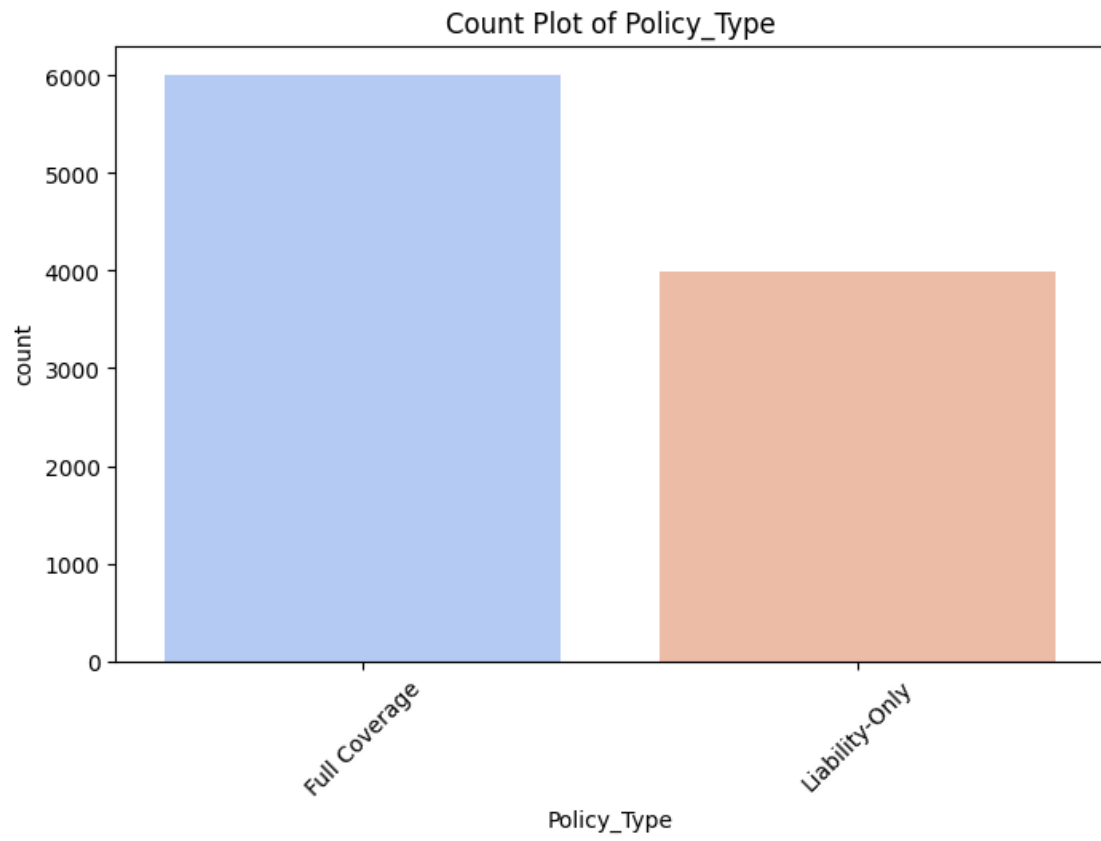


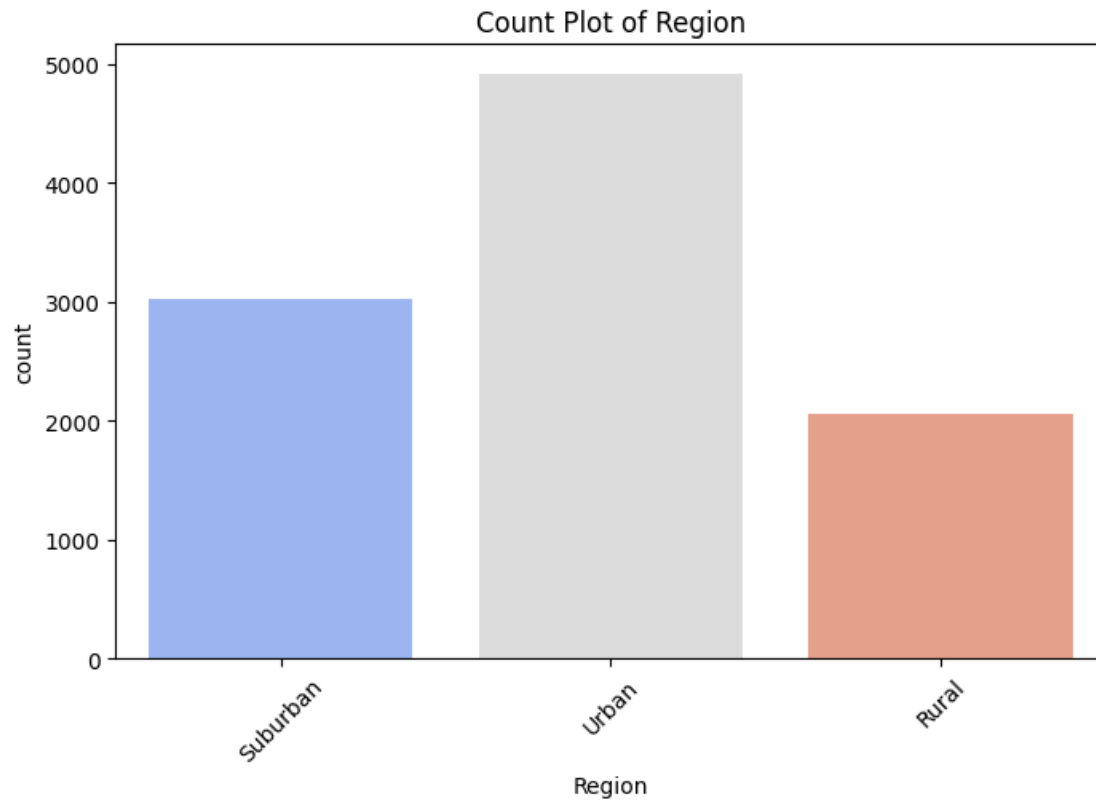


```
[9]: # Count Plots (Categorical Data)
categorical_cols = ['Marital_Status', 'Prior_Insurance', 'Policy_Type', '
↳ 'Region']
for col in categorical_cols:
    plt.figure(figsize=(8, 5))
    sns.countplot(x=df[col], palette='coolwarm')
    plt.title(f'Count Plot of {col}')
    plt.xticks(rotation=45)
    plt.show()
```



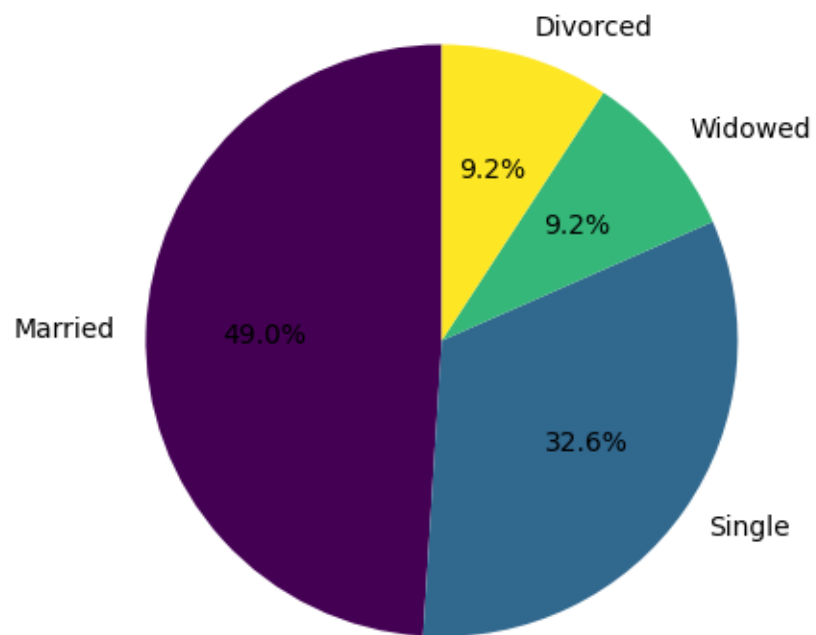







```
[10]: # Pie Charts
for col in categorical_cols:
    plt.figure(figsize=(8, 5))
    df[col].value_counts().plot.pie(autopct='%1.1f%%', startangle=90,
    cmap='viridis')
    plt.title(f'Pie Chart of {col}')
    plt.ylabel('')
    plt.show()
```


Pie Chart of Marital\_Status



Pie Chart of Prior\_Insurance




Pie Chart of Policy\_Type







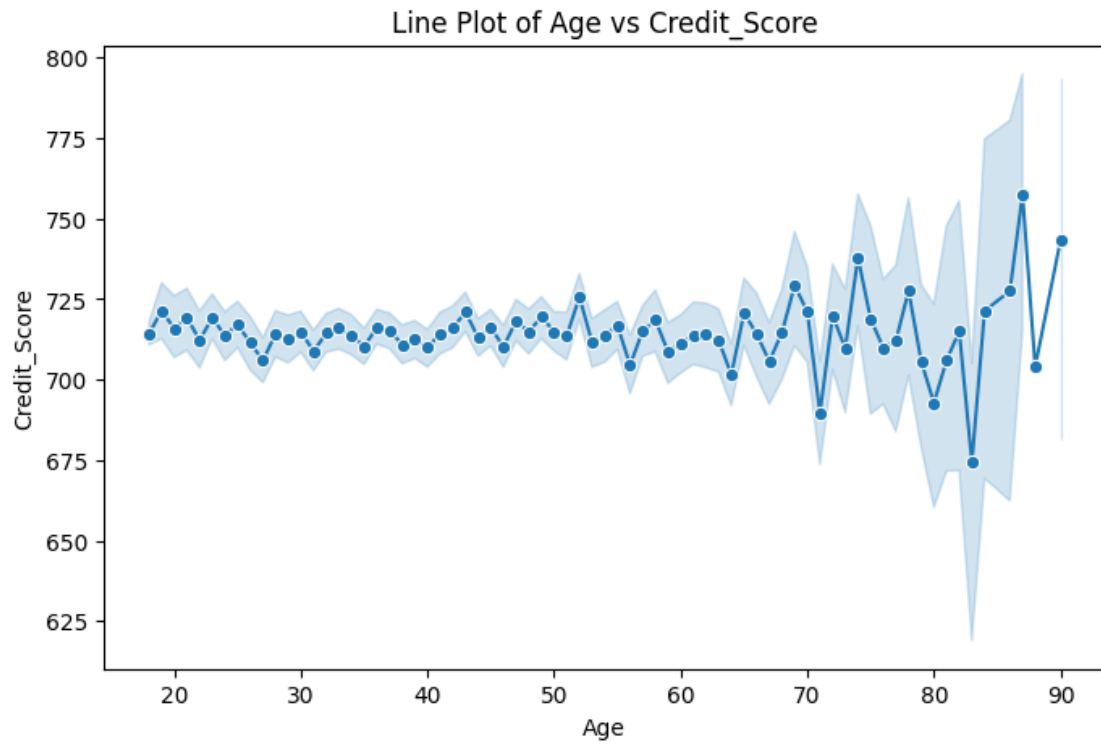
Pie Chart of Region




```
[11]: # Scatter Plots
bivariate_cols = [('Age', 'Credit_Score'), ('Claims_Frequency', 'Claims_Severity')]
for x, y in bivariate_cols:
    plt.figure(figsize=(8, 5))
    sns.scatterplot(x=df[x], y=df[y], alpha=0.6)
    plt.title(f'Scatter Plot of {x} vs {y}')
    plt.show()
```

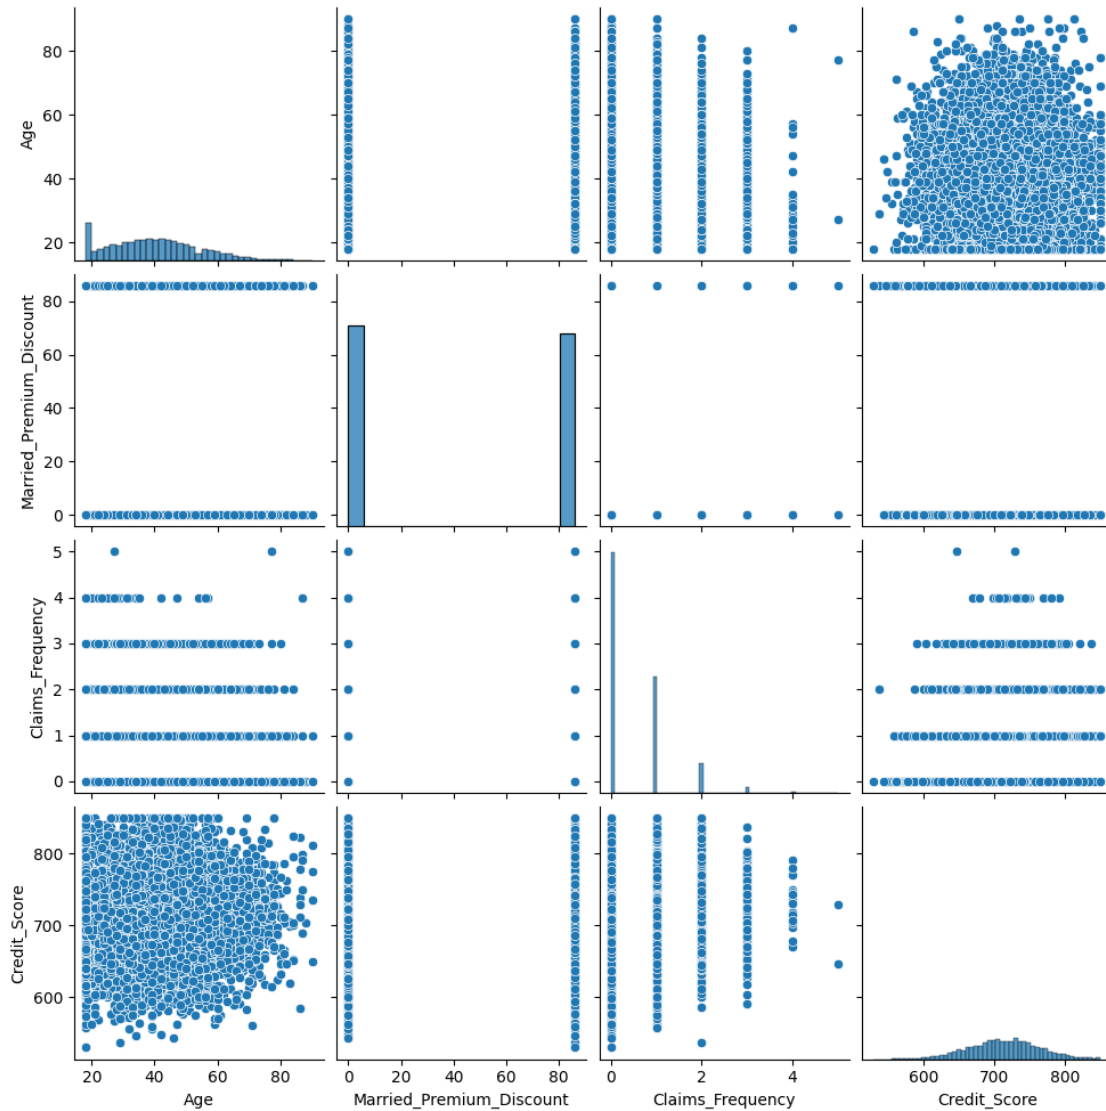


```
[12]: # Line Plots
for x, y in bivariate_cols:
    plt.figure(figsize=(8, 5))
    sns.lineplot(x=df[x], y=df[y], marker='o')
    plt.title(f'Line Plot of {x} vs {y}')
    plt.show()
```





```
[13]: # Pair Plot
sns.pairplot(df[num_cols])
plt.show()
```



```
[14]: # ----- Outlier Detection & Removal
      ↪ ----- #

def detect_outliers(df, cols):
    outliers = {}
    for col in cols:
        Q1 = df[col].quantile(0.25)
        Q3 = df[col].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        outliers[col] = df[(df[col] < lower_bound) | (df[col] >
      ↪ upper_bound)][col]
```

```

    return outliers

# Numerical columns to check for outliers
num_cols = ['Age', 'Married_Premium_Discount', 'Claims_Frequency',
            ↪ 'Credit_Score']
outliers = detect_outliers(df, num_cols)
print("Detected Outliers:", outliers)

# Remove outliers
def remove_outliers(df, cols):
    for col in cols:
        Q1 = df[col].quantile(0.25)
        Q3 = df[col].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        df = df[(df[col] >= lower_bound) & (df[col] <= upper_bound)]
    return df

df = remove_outliers(df, num_cols)
print("Outliers removed. New dataset shape:", df.shape)

```

Detected Outliers: {'Age': 209      90

478	86
1615	87
1957	87
2305	86
2506	84
2521	84
2895	90
3241	83
3716	88
3982	87
4047	82
4870	86
4997	86
5157	83
5224	84
5673	82
5796	86
5846	83
6738	84
6891	90
7223	86
7330	84
7872	82
7874	82

```

8165    84
8248    90
8726    87
Name: Age, dtype: int64, 'Married_Premium_Discount': Series([], Name:
Married_Premium_Discount, dtype: int64), 'Claims_Frequency': 68      3
260     4
284     3
392     4
452     3
..
9868    3
9879    3
9887    3
9890    3
9965    4
Name: Claims_Frequency, Length: 164, dtype: int64, 'Credit_Score': 42      850
94      850
239     850
328     577
384     574
...
9498    850
9728    576
9805    562
9871    577
9937    850
Name: Credit_Score, Length: 81, dtype: int64}
Outliers removed. New dataset shape: (9728, 27)

```

```

[25]: # Data preprocessing
le = LabelEncoder()
categorical_cols = ['Marital_Status', 'Prior_Insurance', 'Claims_Severity',
                    'Policy_Type', 'Region', 'Source_of_Lead']
for col in categorical_cols:
    df[col] = le.fit_transform(df[col])

X = df.drop(columns=['Conversion_Status']) # Target variable
y = df['Conversion_Status']

```

```

[24]: df['Source_of_Lead']

```

```

[24]: 0      Agent
      1      Online
      2      Online
      3      Online
      4      Agent
      ...

```

```

9995    Online
9996    Agent
9997    Agent
9998    Agent
9999    Agent
Name: Source_of_Lead, Length: 9728, dtype: object

```

```

[27]: # Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

# Standardization
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

```

```

[28]: # Define models
models = {
    'Logistic Regression': LogisticRegression(),
    'KNN': KNeighborsClassifier(),
    'Decision Tree': DecisionTreeClassifier(),
    'Random Forest': RandomForestClassifier(),
    'SVM': SVC(probability=True),
    'Naive Bayes': GaussianNB(),
    'Gradient Boosting': GradientBoostingClassifier(),
    'XGBoost': XGBClassifier(),
    'LightGBM': LGBMClassifier(),
    'CatBoost': CatBoostClassifier(verbose=0)
}

```

```

[29]: # Train and evaluate models
results = {}
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    auc = roc_auc_score(y_test, model.predict_proba(X_test)[: , 1])
    results[name] = {'Accuracy': acc, 'AUC': auc}
    print(f"{name} Classification Report:\n", classification_report(y_test,
    ↪y_pred))
    print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

```

Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	817
1	1.00	1.00	1.00	1129



accuracy			1.00	1946
macro avg	1.00	1.00	1.00	1946
weighted avg	1.00	1.00	1.00	1946

Confusion Matrix:

```
[[ 817   0]
 [   0 1129]]
```

KNN Classification Report:

	precision	recall	f1-score	support
0	0.98	0.93	0.96	817
1	0.95	0.99	0.97	1129

accuracy			0.96	1946
macro avg	0.97	0.96	0.96	1946
weighted avg	0.96	0.96	0.96	1946

Confusion Matrix:

```
[[ 762   55]
 [   15 1114]]
```

Decision Tree Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	817
1	1.00	1.00	1.00	1129

accuracy			1.00	1946
macro avg	1.00	1.00	1.00	1946
weighted avg	1.00	1.00	1.00	1946

Confusion Matrix:

```
[[ 817   0]
 [   0 1129]]
```

Random Forest Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	817
1	1.00	1.00	1.00	1129

accuracy			1.00	1946
macro avg	1.00	1.00	1.00	1946
weighted avg	1.00	1.00	1.00	1946

Confusion Matrix:

```
[[ 817   0]
 [   0 1129]]
```

SVM Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	817
1	1.00	1.00	1.00	1129
accuracy			1.00	1946
macro avg	1.00	1.00	1.00	1946
weighted avg	1.00	1.00	1.00	1946

Confusion Matrix:

```
[[ 817   0]
 [   0 1129]]
```

Naive Bayes Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	817
1	1.00	1.00	1.00	1129
accuracy			1.00	1946
macro avg	1.00	1.00	1.00	1946
weighted avg	1.00	1.00	1.00	1946

Confusion Matrix:

```
[[ 817   0]
 [   0 1129]]
```

Gradient Boosting Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	817
1	1.00	1.00	1.00	1129
accuracy			1.00	1946
macro avg	1.00	1.00	1.00	1946
weighted avg	1.00	1.00	1.00	1946

Confusion Matrix:

```
[[ 817   0]
 [   0 1129]]
```

XGBoost Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	817
1	1.00	1.00	1.00	1129
accuracy			1.00	1946
macro avg	1.00	1.00	1.00	1946
weighted avg	1.00	1.00	1.00	1946

Confusion Matrix:

```
[[ 817    0]
 [   0 1129]]
```

[LightGBM] [Info] Number of positive: 4490, number of negative: 3292

[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.002418 seconds.

You can set `force\_row\_wise=true` to remove the overhead.

And if memory is not enough, you can set `force\_col\_wise=true`.

[LightGBM] [Info] Total Bins 485

[LightGBM] [Info] Number of data points in the train set: 7782, number of used features: 26

[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.576973 -> initscore=0.310357

[LightGBM] [Info] Start training from score 0.310357

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[illegible]

```

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

```

LightGBM Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	817
1	1.00	1.00	1.00	1129
accuracy			1.00	1946
macro avg	1.00	1.00	1.00	1946
weighted avg	1.00	1.00	1.00	1946

Confusion Matrix:

```

[[ 817   0]
 [   0 1129]]

```

CatBoost Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	817
1	1.00	1.00	1.00	1129
accuracy			1.00	1946
macro avg	1.00	1.00	1.00	1946
weighted avg	1.00	1.00	1.00	1946

Confusion Matrix:

```

[[ 817   0]
 [   0 1129]]

```


```

[30]: # Model comparison
results_df = pd.DataFrame(results).T
print(results_df.sort_values(by='AUC', ascending=False))

```

	Accuracy	AUC
Logistic Regression	1.000000	1.000000
Decision Tree	1.000000	1.000000
Gradient Boosting	1.000000	1.000000
Random Forest	1.000000	1.000000
SVM	1.000000	1.000000
Naive Bayes	1.000000	1.000000
LightGBM	1.000000	1.000000
XGBoost	1.000000	1.000000
CatBoost	1.000000	1.000000
KNN	0.964029	0.989762

```
[31]: # Feature importance (Random Forest example)
rfc = RandomForestClassifier()
rfc.fit(X_train, y_train)
feature_importances = pd.Series(rfc.feature_importances_, index=df.
    ↪drop(columns=['Conversion_Status']).columns)
feature_importances.nlargest(10).plot(kind='barh')
plt.title("Top 10 Feature Importances")
plt.show()
```



## 0.7 About the Author

Name: Arif Mia

**Profession:** Machine Learning Engineer & Data Scientist

---

### 0.7.1 Career Objective

My goal is to contribute to groundbreaking advancements in artificial intelligence and data science, empowering companies and individuals with data-driven solutions. I strive to simplify complex challenges, craft innovative projects, and pave the way for a smarter and more connected future.

As a **Machine Learning Engineer** and **Data Scientist**, I am passionate about using machine learning, deep learning, computer vision, and advanced analytics to solve real-world problems. My expertise lies in delivering impactful solutions by leveraging cutting-edge technologies.

---

### 0.7.2 Skills

- **Artificial Intelligence & Machine Learning**
  - **Computer Vision & Predictive Analytics**
  - **Deep Learning & Natural Language Processing (NLP)**
  - **Python Programming & Automation**
  - **Data Visualization & Analysis**
  - **End-to-End Model Development & Deployment**
- 

### 0.7.3 Featured Projects

#### **Lung Cancer Prediction with Deep Learning**

Achieved 99% accuracy in a computer vision project using 12,000 medical images across three classes. This project involved data preprocessing, visualization, and model training to detect cancer effectively.

#### **Ghana Crop Disease Detection Challenge**

Developed a model using annotated images to identify crop diseases with bounding boxes, addressing real-world agricultural challenges and disease mitigation.

#### **Global Plastic Waste Analysis**

Utilized GeoPandas, Matplotlib, and machine learning models like RandomForestClassifier and CatBoostClassifier to analyze trends in plastic waste management.

#### **Twitter Emotion Classification**

Performed exploratory data analysis and built a hybrid machine learning model to classify Twitter sentiments, leveraging text data preprocessing and visualization techniques.

---

#### 0.7.4 Technical Skills

- **Programming Languages:** Python , SQL , R
  - **Data Visualization Tools:** Matplotlib , Seaborn , Tableau , Power BI
  - **Machine Learning & Deep Learning:** Scikit-learn , TensorFlow , PyTorch
  - **Big Data Technologies:** Hadoop , Spark
  - **Model Deployment:** Flask , FastAPI , Docker
- 

#### 0.7.5 Connect with Me

**Email:** arifmiahcse@gmail.com

**LinkedIn:** [www.linkedin.com/in/arif-miah-8751bb217](https://www.linkedin.com/in/arif-miah-8751bb217)

**GitHub:** <https://github.com/Arif-miad>

**Kaggle:** <https://www.kaggle.com/arifmia>

Let's turn ideas into reality! If you're looking for innovative solutions or need collaboration on exciting projects, feel free to reach out.

---

How does this look? Feel free to suggest changes or updates!